

**TECHGUARD · UNIFIED SERVICE PLATFORM**

Walkthrough Deck · June 2026

BUILT ON SALESFORCE AGENTFORCE**• AGENTIC AI · IT SERVICE RE-PLATFORM**

TechGuard **USP**

The agentic operating system for managed IT services.

Re-platforming a 3,000-employee MSP across 2M endpoints, 15K B2B clients, and three personas in 36 months — on Salesforce Agentforce, Data Cloud, Slack, and MuleSoft.

KAUSTUBH CHAUDHARY

Companion artifacts: Product Strategy & Specification Document · Interactive Prototype · Architecture Diagram

2M

ENDPOINTS

15K

B2B CLIENTS

36mo

HORIZON

01 / 19

THE CONTRARIAN POV

The ticket is dying.

Most ITSM vendors are racing to bolt AI onto the ticket. They're building faster horses.

2022 — 2025

01

Augmented reactive

Chatbots, copilots, summarization. The ticket survives; humans get faster. **Most enterprises are still here.**

"AI as productivity feature"

• WE ARE HERE

02

2025 — 2027

Conversational autonomous

Agents resolve top-volume cases end-to-end. **Salesforce: 96% auto-resolution at 100K employees. McAfee replaced ServiceNow.**

"The ticket becomes an audit log"

2027 +

03

Proactive, agent-initiated

Agents act *before* the customer notices. Telemetry triggers anomaly detection. **Service becomes a trust-building surface.**

"The category itself moves"

⚡ Unit of competitiveness

shifts from ticket-throughput to **outcome-velocity**

🏢 L1/L2/L3 org chart collapses

for the 60–80% of patternable work

🌐 Data substrate is the moat

tools commoditize; **data + action does not**

🛡️ Governance is infrastructure

audit, explainability, blast-radius controls

For TechGuard, the strategic question isn't *"when do we add AI features."* It is: **do we re-platform before the category does it to us.**

THE STARTING STATE

TechGuard runs manual, siloed, reactive — at scale.

3,000 employees · 15K B2B clients · 2M endpoints across CCTV, Cloud, Connectivity. The current paradigm is the cost structure we are dissolving.

Reactive support · clients call when broken
High MTTR · SLA penalties · emergency truck rolls

Siloed teams across 3 service lines
Same client = 3 tickets for 1 root cause · cross-domain incidents handled poorly

Manual triage · CSAs swivel 4-5 tools
30-40% of CSA time on triage · high attrition

Truck-roll waste · wrong parts / no fault found
20-30% of rolls wasted · high CAC of service delivery

Knowledge fragmented · 60-90 day CSA ramp
KB stale or contradictory · expertise walks out

No predictive layer · telemetry untapped
Predictable failures become P1 incidents

QUANTIFIED COST-OUT · ANNUAL

Frontline service cost

2.4M contacts × \$18 fully-loaded

\$43M

Truck-roll waste

25% × 200K rolls × \$450/roll

\$22M

SLA penalty exposure

3-7% of contract value

\$10-25M

TOTAL ADDRESSABLE COST-OUT

\$75-90M

/ year · before any revenue upside

THE VISION · ONE SENTENCE

USP becomes the agentic operating system for managed IT services — the platform on which every TechGuard service outcome is detected, decided, dispatched, and documented autonomously.

HORIZON 1 · YEAR 1

Augmented Service

AI in the loop · humans close it

- ▶ Endpoint Graph (Data Cloud)
- ▶ Frontline Copilot · Maya
- ▶ Backend Triage Agent · Rohan
- ▶ Field Mobile Companion · Maria
- ▶ Top-20 Use Case Library
- ▶ Trust & Audit Layer v1

Y1 EXIT

50% auto-res · 50 lighthouse customers · Dreamforce '26 reference

HORIZON 2 · YEAR 2

Autonomous Service

From answer to action

- ▶ Autonomous Action Library (100+)
- ▶ Cross-Domain Root Cause Agent
- ▶ Predictive Maintenance v1
- ▶ Per-Client Trust Dial
- ▶ Agentforce Voice + Auto-Dispatch
- ▶ Blast-Radius Classifier

Y2 EXIT

65% auto-res · 25% proactive · ITSM MQ Visionary

HORIZON 3 · YEARS 3–5

Outcome Platform

Reframe the category

- ▶ Outcome-Based Pricing Engine
- ▶ Federated Cross-Tenant Learning
- ▶ Self-Healing Closed Loop
- ▶ Agentic Pre-Sales / Site Survey
- ▶ Vertical Agent Marketplace
- ▶ USP-as-a-Platform (MSP white-label)

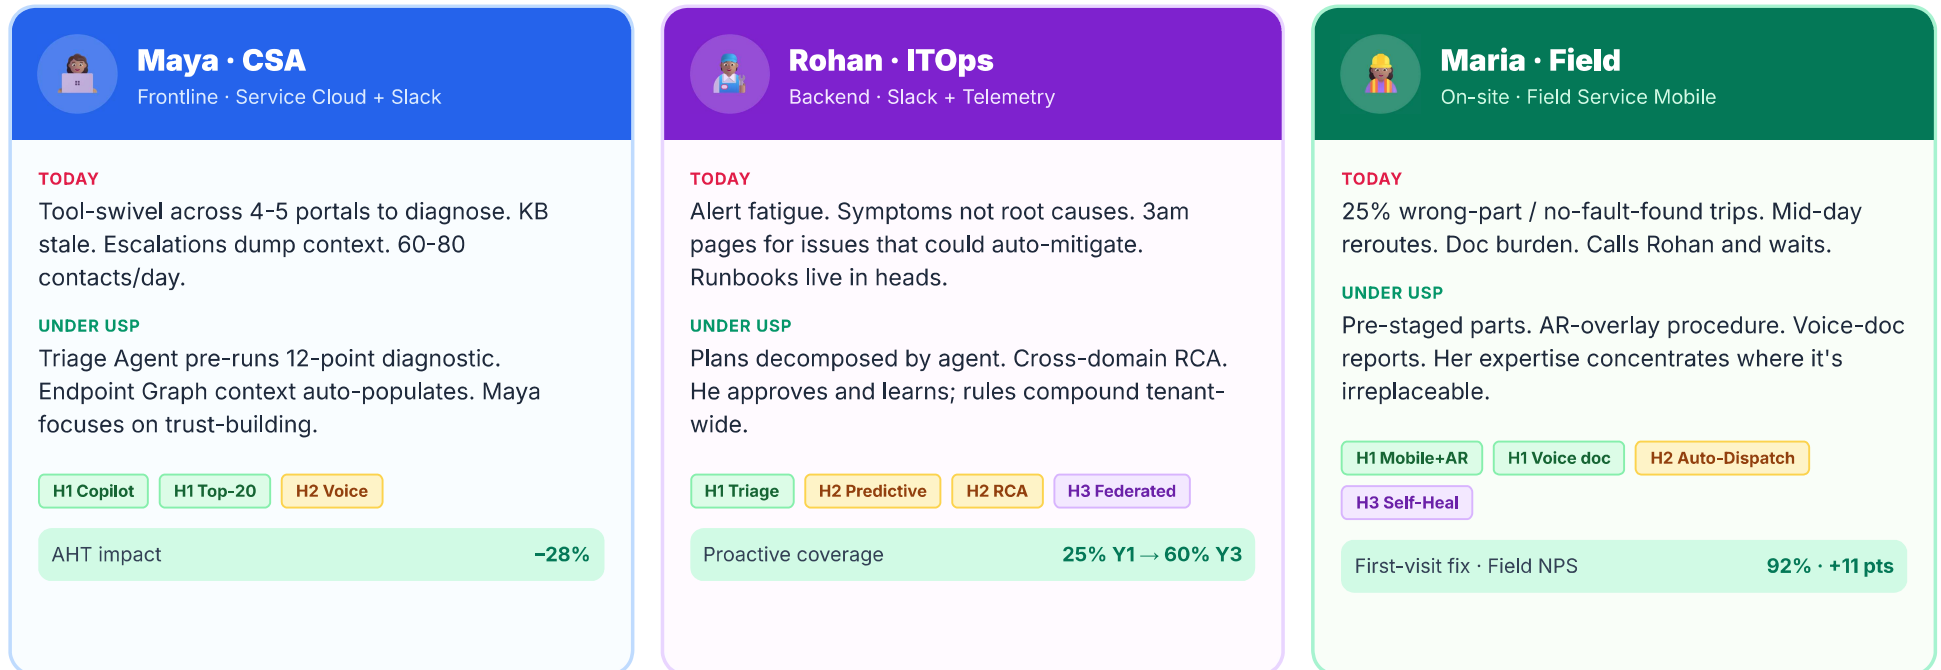
Y3-5 EXIT STATE

80% auto-res · 60% proactive · TechGuard becomes platform vendor

THE UNIFYING PATTERN

Context loss at every handoff *is* the cost structure.

Each handoff today destroys 30–50% of the diagnostic work done before it. Today's L1/L2/L3 architecture isn't a bug — it's the embodiment of that loss. **Solve the handoffs and you collapse the cost structure.**



COMPOUNDING-WEIGHTED SCORING

The Endpoint Graph isn't the loudest pain. It's the most *compounding* one.

Score = (Frequency × Severity × **Strategic Lift**) ÷ Implementation Effort. Strategic Lift is the multiplier — solving the substrate unlocks 8–10 other solutions architecturally.

PAIN	PERSONA	FREQ	SEV	STRAT LIFT	IMPL EFFORT	SCORE	SHIP
Fragmented endpoint data across CCTV/Cloud/Conn.	All	5	5	★ 5	4	31.3	H1
Tool-swivel during diagnosis	Maya	5	4	4	2	40.0	H1
Alert fatigue · symptoms not root causes	Rohan	5	5	4	3	33.3	H1→H2
Wrong-part / no-fault-found truck rolls	Maria	4	5	3	3	20.0	H1
Reactive vs proactive posture	All	4	5	★ 5	4	25.0	H2
Cross-domain incidents fall between teams	All	3	5	★ 5	4	18.8	H2
Documentation burden post-visit	Maria	5	3	2	1	30.0	H1
No closed-loop learning from resolved cases	Rohan	3	4	★ 5	3	20.0	H2→H3

THE INSIGHT

Four pains share **Strategic Lift = 5**. Three of them are unbuildable without the fourth — **the Endpoint Graph**. That makes it the Y1 anchor bet, not the most visible.

IMPLICATION FOR SEQUENCING

Ship Endpoint Graph **before** any agent capability that depends on it. Building Frontline Copilot on a stub data layer creates re-platforming debt within 6 months.

FAST WIN FOR H1

Maria's voice-doc shipped to all Field Techs in 30 days · daily frequency · low complexity · trust-building win that buys runway for the bigger bets.

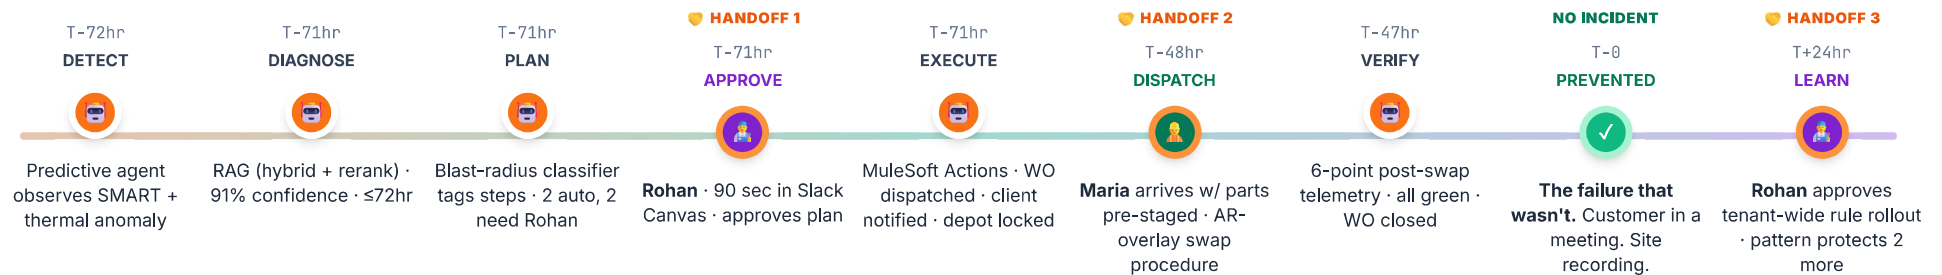
PART 2 · CUSTOMER JOURNEY MAP

TECHGUARD USP · K. CHAUDHARY

THE CANONICAL JOURNEY

Predictive server failure · 3 personas, 3 handoffs, zero incident.

Same incident travels across the platform. The disk that would have hard-failed at T-0 — never did.

**\$42K**

SLA PENALTY AVOIDED

1

PREVENTIVE TRUCK ROLL (NOT EMERGENCY)

0

CSA TICKETS · CUSTOMER NEVER KNEW

18 SOLUTIONS · PRIORITIZED · SEQUENCED

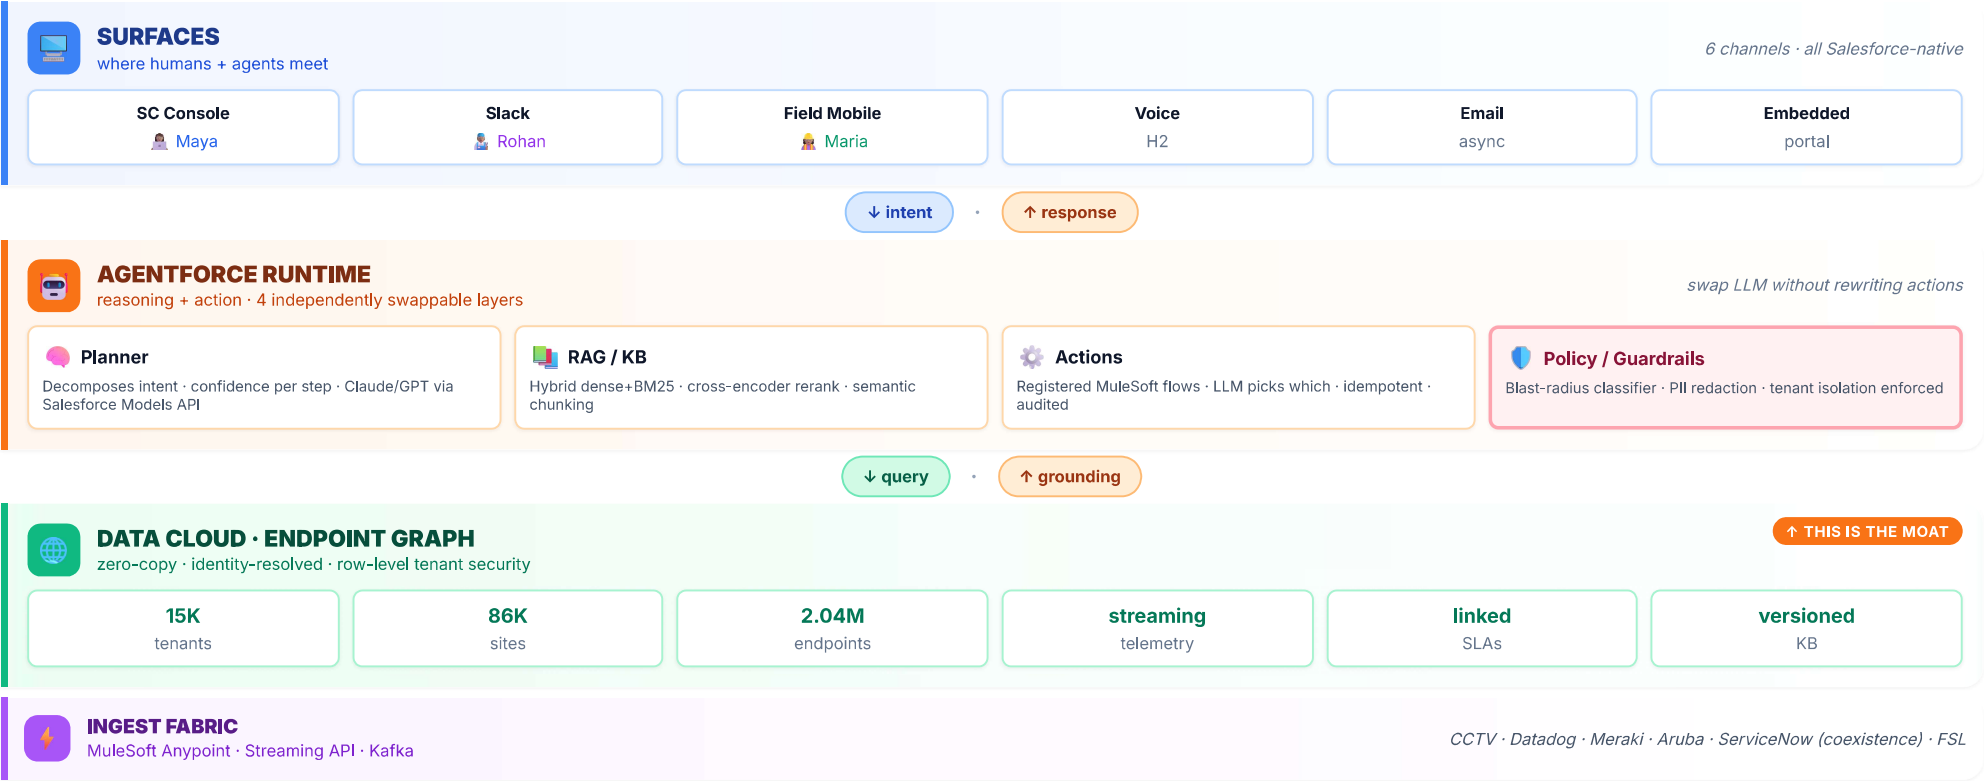
Substrate first. Action library second. Predictive only when the data is trusted.

H1 · YEAR 1 BUILD ORDER H1 S1 → (S2, S3, S4 II) → S5 → S6 S1 Endpoint Graph (Data Cloud) 35.0 S2 Frontline Copilot 32.0 S3 Backend Triage + Alert Correlation 28.5 S4 Field Mobile Companion + AR + Voice 26.0 S5 Top-20 Use Case Library 22.0 S6 Trust & Audit Layer ⚙️	H2 · YEAR 2 BUILD ORDER H2 S7 → S8 → (S9, S10, S11, S12 II) S7 Autonomous Action Library + Blast-Radius 24.0 S8 Predictive Maintenance v1 22.5 S9 Cross-Domain Root Cause Agent 20.0 S10 Per-Client Trust Dial 18.0 S11 Agentforce Voice (CSA) 17.0 S12 Field Auto-Dispatch + Route Opt 16.0	H3 · YEARS 3-5 BUILD ORDER H3 S13 + S16 → S14, S15, S17 → S18 S13 Outcome-Based Pricing Engine ★ S14 Federated Cross-Tenant Learning 18.0 S15 Agentic Pre-Sales / Site Survey 14.0 S16 Self-Healing Closed Loop (Tier-0) ★ S17 Vertical Agent Marketplace 12.0 S18 USP-as-a-Platform (MSP white-label) ★
--	--	--

THE NON-DO LIST (A DIRECTOR-LEVEL DISCIPLINE)			
✗ No standalone client portal customers consume via Slack / email / voice / embedded	✗ No ServiceNow feature-parity race change mgmt, CMDB are coexistence plays	✗ No bespoke per-vertical features in H1 marketplace play is H3, not now	✗ No fine-tuned models in H1 RAG + prompt eng covers 90%

THREE LAYERS · ONE SUBSTRATE

Surfaces · Agentforce Runtime · Data Cloud Endpoint Graph.



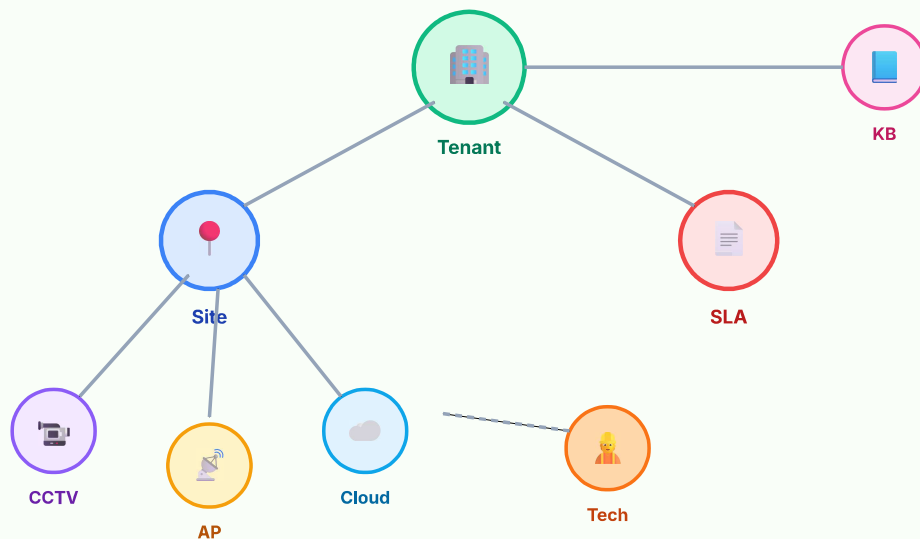
PRINCIPLE Planner creative · Actions deterministic · Policy independent.

→ Swap any layer without rewriting the others.

THE ENDPOINT GRAPH IS THE MOAT

Identity-resolved unified data is a 6-month build for us. A multi-year one for ServiceNow.

SCHEMA · ONE GRAPH SPANNING ALL DATA



WHY THIS BEATS INCUMBENTS

- vs ServiceNow** CMDB-bound. Cannot easily span CCTV / Cloud / Connectivity in one identity-resolved graph.
- vs Workday Sana** HR-bound. Has no equivalent telemetry-substrate for endpoint operations.
- vs Microsoft Copilot** Tightly coupled to specific apps. Cross-system orchestration is weak.

Five data principles

- ▶ **Single source of truth** = Data Cloud · no parallel data lake
- ▶ **Tenant isolation** at the data layer · not the application layer
- ▶ **Telemetry is first-class** · co-located with customer records
- ▶ **Knowledge versioned + attributable** · with freshness scores
- ▶ **CCTV video never ingested** · only metadata · GDPR/CCPA safe

Vendor lock-in mitigated

Data Cloud is Iceberg-compatible. The Endpoint Graph schema is portable. If we ever needed to exit Agentforce, the data substrate exits with us.

SEPARATION OF CONCERNS

Planner creative. Actions deterministic. Policy independent.

The single most important architectural decision. It is what prevents the LLM from inventing a `cancelPayroll` tool that doesn't exist.



LAYER 1 Planner

Decomposes intent into a parallelizable plan with confidence scores per step.

► **LLM swappable**

Claude Opus 4.7 · GPT-5 · Gemini 2.5 via Salesforce Models API.

Per-use-case routing trades off cost · latency · quality.



LAYER 2 RAG / KB

Hybrid retrieval: dense embeddings + BM25 keyword.

► **Cross-encoder rerank**

Semantic chunking, not fixed-window. KB articles versioned with freshness scores.

Hybrid catches alphanumerics (SKUs, IPs) that pure semantic search misses.



LAYER 3 Actions

Each action = a registered MuleSoft flow.

► **LLM picks WHICH, not how**

Actions are deterministic. The LLM can never invent a tool that doesn't exist. Every call audit-logged.

100+ Actions by H2 · idempotent · reversible where possible.



LAYER 4 Policy

Owns safety bar *independent* of model upgrades.

► **Blast-radius classifier**

LOW · MEDIUM · HIGH tagging per action. Per-tenant Trust Dial governs threshold. PII redaction. Kill switch.

Tenant isolation enforced here · zero cross-tenant leaks SLA.

⊘ **Hallucinated diagnosis**

Retrieval grounding required for factual claims. Below threshold → escalate.

⊘ **Action misfire**

Blast-radius classifier gates writes. HIGH = approval required.

⊘ **Cross-tenant leak**

Tenant ID required parameter on every retrieval. Row-level security.

⊘ **Model drift**

500-test golden eval suite per use case. Automated regression block.

TRUST IS THE RATE LIMIT ON AUTONOMY

Autonomy without trust isn't speed. It's **blast radius**.

CUSTOMER-OWNED TRUST DIAL

The autonomy threshold is set by the customer, not by us.

LOW BLAST

Restarts · queries · diagnostics · snapshots

AUTO

MEDIUM BLAST

Firmware push · config change · work order dispatch

ITOPS APPROVAL

HIGH BLAST

Mass change · financial · termination · cross-tenant

CLIENT APPROVAL

- Default conservative · opt-in to higher autonomy
- Reviewed quarterly per client · tied to SLA terms
- Every threshold change audit-logged · explainable

HARD GUARDRAILS · CANNOT BE IMPROVED AT THE COST OF

CSAT	never -2 pts vs baseline
Hallucination rate	< 1% top-20
Escalation accuracy	> 90%
PII / data leak	zero · hard kill
Cross-tenant leak	zero · row-level security
Agent error → client P1	< 1 / quarter

EVAL-AS-CI · NON-NEGOTIABLE

- 500 golden test cases per top use case
- Automated regression block before any model / prompt change ships
- Per-tenant accuracy dashboards in customer console

"No agent ships without an eval suite."

FUNCTIONAL PROTOTYPE · CLICKABLE

One platform. Three surfaces. The journey, end-to-end.

Single unified prototype — not three separate persona apps — because USP is unified by definition.

home

AI · IT RE-PLATFORM

The agentic OS for managed IT services.

▶ Canonical demo

87.4
AROE

62%
Auto

27%
Proact

+9
NPS

HOME · thesis · NSM · 3 personas · roadmap

canonical-demo

T~71hr · HANDOFF #1

Agent → Rohan (ITOps)

Predictive Maintenance Agent

⚠ NVR-03 · Site 4 · 91% · ≤72hr

Plan ready (4 steps).

Approve

Edit

Reject

9 steps · 3 handoffs · arrow keys

★ CANONICAL DEMO · predictive failure E2E

metrics-architecture

SURFACES · 3 personas · 6 channels

AGENTFORCE · Planner · RAG · Actions · Policy

DATA CLOUD · ENDPOINT GRAPH

INGEST · MuleSoft · Kafka

Source systems · CCTV · Cloud · Connectivity

\$65-70M Y3 · ~16mo payback

METRICS · ARCHITECTURE · UNIT ECONOMICS

Maya · CSA — Service Cloud Console mock with Endpoint Graph + Agent Assist

Rohan · ITOps — Slack-style alerts, RCA, eval/audit sidebar

Maria · Field — Mobile + AR-overlay + voice doc

LIVE PROTOTYPE URL

https://kaustubh-vibecoder.github.io/techguard-usp-prototype/

Path

Home → ▶ Canonical Demo → 3 personas → Metrics

13 / 19

file:///C:/Users/kchaudhary/Claude Docs/PM-Interview-Prep/Salesforce Interview/salesforce-takehome-techguard-deck.html

13/19

NORTH STAR METRIC

AROE — Autonomous Resolved Outcomes per 1K Endpoints / month.

Service outcomes (resolutions, preventive interventions, requests fulfilled) **fully executed without human touch** AND **verified by post-resolution telemetry or client confirmation**, normalized to 1,000 endpoints.

CURRENT VS TARGETS

87.4

today
(50 lighthouse cohort)

100

Y1 target

~100 / 80% auto

Y3 steady state

WHY THIS NSM AND NOT OTHER CANDIDATES

- ✗ **Deflection rate** — gameable; high deflection at cost of unresolved escalations is a worse outcome
- ✗ **MTTR** — doesn't capture proactive work that never becomes an incident
- ✗ **CSAT / NPS** — lagging; many non-product factors
- ✗ **Tickets closed** — the wrong unit of work; anchors to the paradigm we're dismantling

SUPPORTING GOAL TREE · 6 INPUT METRICS

G1

Proactive Coverage

% incidents prevented before client awareness

27% Y1 target 25% → Y3: 60%

G2

Auto Resolution Rate

% contacts resolved without human touch

62% Y1: 50% → Y3: 80%

G3

First-Visit Fix Rate

% truck rolls resolving first visit (Field)

84% Y1: 80% → Y3: 92%

G4

MTTR

Detection → verified resolution (min)

-43% Y1: -40% → Y3: -65%

G5

Service Cost / Endpoint

Fully-loaded annual cost per endpoint

-21% Y1: -20% → Y3: -45%

G6 · GUARDRAIL

Client Trust Index

Composite CSAT + NPS + SLA + survey

+9 NPS pts · never below -2

COUNTER-METRICS CATCH WHAT NSM CAN'T

If AROE goes up while Maya stops growing or Maria starts hating her job — we broke it.

Counter-metrics are surfaced on equal footing with NSM in the metrics dashboard. They are the difference between **scaling leverage** and **scaling exploitation**.

 **GUARDRAILS · THE FLOOR**

CSAT
Never below baseline by >2 pts

4.7 / 5

Hallucination rate
< 1% on top-20 use cases

0.4%

Escalation accuracy
> 90% (escalations actually need humans)

94%

PII leaks · cross-tenant
Zero. Hard kill switch. Board notification.

0 · 0

Agent error → client P1
< 1 / quarter

0

 **COUNTER-METRICS · CATCH THE GAMING**

Re-open rate
Catches "false resolution" — agent marks closed, client re-opens.

2.1% ↓

Esc-after-deflect
Catches "deflect at all costs" — claims resolution, customer escalates <48hr.

3.4%

CSAT delta · agent vs human
If agent CSAT diverges down, deflection winning over trust.

+0.1

AWU / resolved outcome
Catches token-cost runaway. Protects gross margin.

\$0.94

Field Tech NPS
"We didn't automate Maria's dignity away."

+11 pts

CSA career mobility
If Maya doesn't progress faster, we built leverage on her back.

3.2× promotion rate

15 / 19

THE MATH, NOT THE DIRECTIONAL

\$65–70M annual P&L impact at Y3. ~16-month payback. The numbers compose.

Y3 UNIT ECONOMICS · WORKING FROM VOLUME TO P&L

Inbound contacts (2M endpoints × 1.2/yr) + 0.6M proactive interventions	3.0M / yr
Autonomous resolution (80% Y3 target) = AROE ≈ 100	2.4M / yr
Flex Credit COGS per outcome (200 credits avg) 200 credits × \$0.005 = \$1.00 / outcome	(\$2.4M)
Direct contact-cost displaced 2.4M × \$18 fully-loaded	+\$43M
Truck-roll waste reduction 40% reduction · ~50K rolls avoided	+\$15M
Avoided SLA penalty exposure Predictive layer · prevented P1s	+\$10–15M
Net annual P&L impact (Y3)	~\$65–70M

Y1 INVESTMENT ASK

\$24–27M

1 Group PM + 5 PMs	\$3.0–3.5M
~55 engineers	\$13–15M
4 designers + 2 UXR	\$1.5M
5 CSMs + lighthouse program	\$2.5M
Flex Credits PreCommit	\$3M
Eval infra	\$1M

Payback: ~16 months

CO-ROADMAP UPSIDE FOR SALESFORCE

- ▶ **\$5–8M / yr** Salesforce ARR from TechGuard at steady state (Flex Credits + Data Cloud + Slack)
- ▶ TechGuard becomes the **reference architecture** Salesforce takes to ~400 MSP TAM
- ▶ Conservatively **~\$500M ARR adjacency** across MSPs at 25% conversion over 5 years

FIVE NAMED TRADEOFFS · THE CHOSEN SIDE · THE MANAGEMENT APPROACH

Strategy is choosing what we explicitly will not optimize for.

	🔥 THE TENSION	✓ THE CHOSEN PATH	🔧 MANAGEMENT APPROACH
1	Autonomy vs Trust More autonomy = bigger deflection, bigger blast radius per error.	Trust Dial owned by the customer Per-tenant auto vs human-approved control. Default conservative; opt-in to more autonomy.	First-class product surface, not config screen. Quarterly client review. Tied to SLA terms.
2	Salesforce leverage vs lock-in Native = 6–12 mo TTV. But couples us to their roadmap and pricing.	Embrace lock-in · mitigate via portable data Endpoint Graph schema Iceberg-portable. Runtime isn't. Honest trade accepted.	Annual platform-risk review. Flex Credits PreCommit. Co-launch partner with Service Cloud / Agentforce leadership.
3	ITSM parity vs category redefinition Match ServiceNow breadth (CMDB, change mgmt) — or change the buyer's question.	Coexist · counter-position on outcomes No CMDB depth investment H1–H2. Sell ROI of outcomes, not seats in ITSM console.	Sales enablement around outcome ROI. SN coexistence reference architecture. Explicit CMDB non-investment.
4	Standardization vs customization 15K clients each want their flow. Custom = maintenance hell; rigid = lost deals.	3-tier model · platform / config / custom T1 standardized · T2 self-service config · T3 custom (Tier-0 only, via PS).	PMs gatekeep T1. Pricing makes T3 self-funding via PS revenue. No T3 in H1.
5	Speed-to-market vs eval rigor Could ship Top-20 in 3 months without eval. Trust incidents would follow.	Eval framework is Day-1 deliverable 4–6 eng weeks upfront on golden tests + regression detection. Before any agent ships.	"No agent ships without an eval suite" — engineering review gate. Trust incidents publicly postmortem'd.

7 THINGS I DON'T YET KNOW

Here's how I'd find out in the first 60 days — before committing the full Y1 investment.

1 Vendor API coverage

Do MuleSoft accelerators exist for Axis, Hikvision, Meraki, Aruba, Datadog?

→ **Vendor-by-vendor API audit, week 1**

2 Existing ServiceNow footprint

Coexistence architecture is more nuanced than displacement if 30%+ clients have SN.

→ **Internal SN audit + top-20 client survey, weeks 2-3**

3 Client autonomy appetite

Trust Dial assumption depends on customers wanting it. If lower than expected, H2 slips.

→ **Survey 50 lighthouse candidates on action categories**

4 Field tech adoption barriers

Union agreements, mobile data plans, depot constraints, AR-device support.

→ **Ride-alongs with 3-5 techs, week 3 · review with Field Ops VP**

5 Contact-cost baseline (\$18 assumption)

Industry benchmark; TechGuard-specific may diverge. Unit economics scale to reality.

→ **CFO partnership weeks 1-2 · lock fully-loaded number**

6 Per-vertical regulatory exposure

Healthcare (HIPAA) · financial (SOX, PCI) · public (FedRAMP) change Trust Layer scope.

→ **Portfolio segmentation by regulatory exposure, month 1**

7 ⚠️ Board-level risk: Salesforce co-roadmap commitments

If Salesforce platform priorities shift such that the Endpoint Graph can't be a Data Cloud first-party schema and has to be a custom extension, TTV extends by ~6 months and the moat narrows.

→ **Want a signed commitment from Service Cloud + Data Cloud GMs that USP is in their joint top-3 priorities for FY27. Validate by month 2.**

CLOSING
THE THESIS

SALESFORCE DIRECTOR PM · TAKE-HOME

We're not building an AI feature.

We're re-platforming a service business.

The platform-level moves in the next 36 months decide whether TechGuard is a category-defining MSP or a casualty of the category's redefinition.

WHAT SUCCESS LOOKS LIKE AT END OF YEAR 1

- ▶ AROE baselined and growing MoM
- ▶ 50%+ autonomous resolution on top-20
- ▶ -28% CSA AHT · saved hours → L2 ramp
- ▶ First "prevented predictive failure" case study published
- ▶ **Zero P1 trust incidents**
- ▶ **Salesforce Dreamforce '26 reference** — Marc names TechGuard from stage

WHAT I'M ASKING FOR

- ▶ Alignment on Horizon-1 scope and the non-do list
- ▶ Sign-off on AROE as NSM + guardrail framework
- ▶ Commitment to **eval-as-Day-1** principle
- ▶ Investment in the **Endpoint Graph (S1)** as foundational H1 bet
- ▶ Endorsement of **customer-owned Trust Dial** principle
- ▶ Signed Salesforce **co-roadmap commitments** on platform dependencies

19 / 19

Thank you. The full document, prototype, and architecture diagram are in the submission folder.